

# SQL Intermediate



# Aggregate Functions

- COUNT, MAX, MIN, SUM, AVG
- Example:

```
SELECT max(mark)
FROM enrolment;
```

```
SELECT avg(mark)
FROM enrolment;
```

```
SELECT min(mark)
FROM enrolment;
```

```
SELECT count(stu_nbr)
FROM enrolment
WHERE mark >= 50;
```

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

**Q1. What will be displayed by the following SQL statement?**

```
SELECT count(*), count(mark)
FROM enrolment;
```

- A. 8, 8
- B. 8, 3
- C. 3, 3
- D. 3, 8

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

## Q2. What will be displayed by the following SQL statement?

SELECT count(\*), count(stu\_nbr), count(distinct stu\_nbr)  
FROM enrolment;

- A. 8, 8, 4
- B. 8, 8, 8
- C. 8, 4, 8
- D. 8, 4, 4

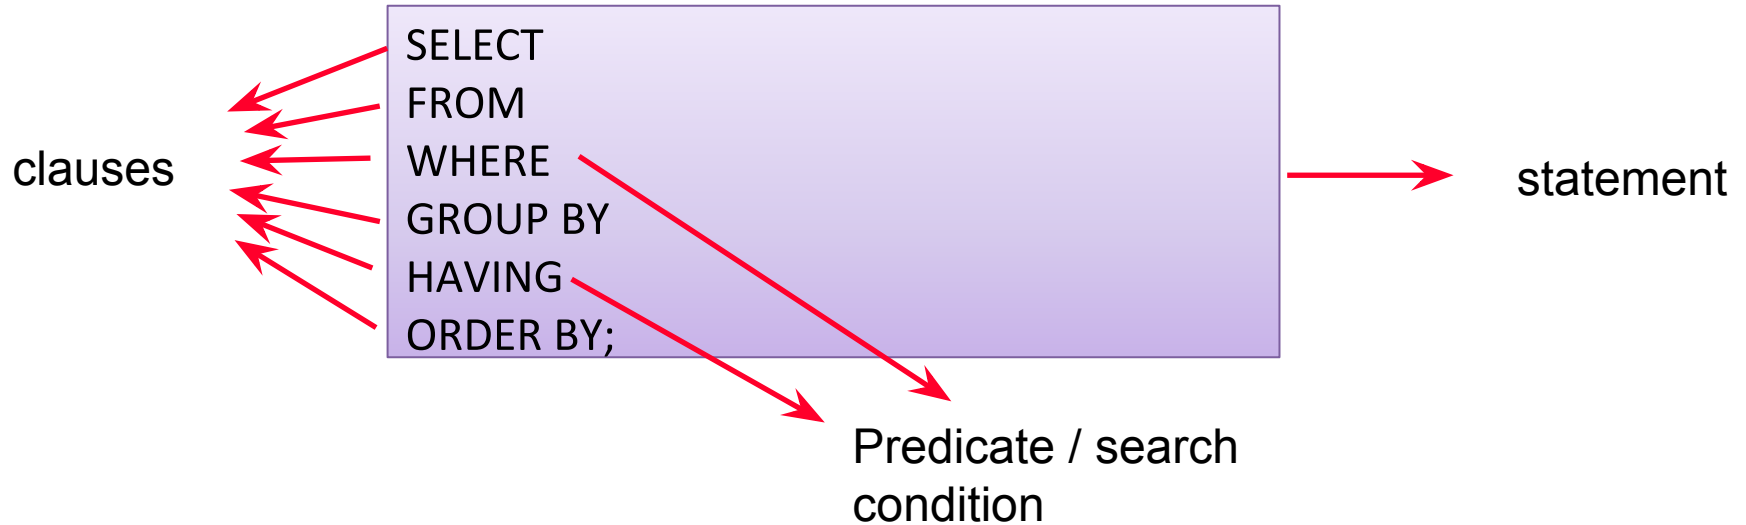
	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

**Q3. We want to calculate the *average mark of the 8 rows* in the above table. What SQL statement should we use?**

**Note: We want to calculate  $(78+35+65)/8=22.25$**

- A. SELECT avg(mark) FROM enrolment;
- B. SELECT sum(mark)/count(mark) FROM enrolment;
- C. SELECT sum(mark)/count(\*) FROM enrolment;
- D. SELECT avg(NVL(mark,0)) FROM enrolment;
- E. None of the above.
- F. More than one option is correct.

# Anatomy of an SQL Statement - Revisited



## GROUP BY

- If a GROUP BY clause is used with aggregate function, the DBMS will apply the aggregate function to the different groups defined in the clause rather than all rows.

```
SELECT avg(mark)  
FROM enrolment;
```

```
SELECT unit_code, avg(mark)  
FROM enrolment  
GROUP BY unit_code  
ORDER BY unit_code;
```

```
SQL>
SQL> SELECT avg(mark)
      2  FROM enrolment;
```

```
      AVG(MARK)
-----
59.3333333
```

```
SQL>
SQL> SELECT unit_code, avg(mark)
      2  FROM enrolment
      3  GROUP BY unit_code
      4  ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)
-----  -----
FIT1001  59.3333333
FIT1002
FIT1004
```



# What output is produced?

```
SELECT avg(mark)
FROM enrolmentA;
```

```
SELECT unit_code, avg(mark)
FROM enrolmentA
GROUP BY unit_code
ORDER BY unit_code;
```

```
SELECT unit_code, avg(mark), count(*)
FROM enrolmentA
GROUP BY unit_code
ORDER BY unit_code;
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SQL> SELECT avg(mark)
      2 FROM enrolmentA;
```

```
AVG(MARK)
-----
      56
```

```
SQL>
SQL> SELECT unit_code, avg(mark)
      2 FROM enrolmentA
      3 GROUP BY unit_code
      4 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)
-----  -
FIT2004      60
FIT2094      50
```

```
SQL>
SQL> SELECT unit_code, avg(mark), count(*)
      2 FROM enrolmentA
      3 GROUP BY unit_code
      4 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----  -
FIT2004      60         3
FIT2094      50         2
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

## What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code, year  
ORDER BY unit_code, year;
```

```
SQL> SELECT unit_code, avg(mark), count(*)
2 FROM enrolmentA
3 GROUP BY unit_code, year
4 ORDER BY unit_code, year;
```

*Note: attributes in the GROUP BY clause do not have to appear in the select list*

UNIT_CO	AVG(MARK)	COUNT(*)
FIT2004	40	2
FIT2004	100	1
FIT2094	20	1
FIT2094	80	1

```
SQL> SELECT unit_code, year, avg(mark), count(*)
2 FROM enrolmentA
3 GROUP BY unit_code, year
4 ORDER BY unit_code, year;
```

UNIT_CO	YEAR	AVG(MARK)	COUNT(*)
FIT2004	2015	40	2
FIT2004	2016	100	1
FIT2094	2015	20	1
FIT2094	2016	80	1

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

## HAVING clause

- It is used to put a condition or conditions on the groups defined by GROUP BY clause.

```
SELECT unit_code, count(*)  
FROM enrolment  
GROUP BY unit_code  
HAVING count(*) > 2;
```

## What output is produced?

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING count(*) > 2  
ORDER BY unit_code;
```

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING avg(mark) > 55  
ORDER BY unit_code;
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SQL> SELECT unit_code, avg(mark), count(*)
2 FROM enrolmentA
3 GROUP BY unit_code
4 HAVING count(*) > 2
5 ORDER BY unit_code;
```

UNIT_CO	AVG(MARK)	COUNT(*)
-----	-----	-----
FIT2004	60	3

```
SQL>
SQL> SELECT unit_code, avg(mark), count(*)
2 FROM enrolmentA
3 GROUP BY unit_code
4 HAVING avg(mark) > 55
5 ORDER BY unit_code;
```

UNIT_CO	AVG(MARK)	COUNT(*)
-----	-----	-----
FIT2004	60	3

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

# HAVING and WHERE clauses

```
SELECT unit_code, count(*)  
FROM enrolment  
WHERE mark IS NULL  
GROUP BY unit_code  
HAVING count(*) > 1;
```

- The WHERE clause is applied to ALL rows in the table.
- The HAVING clause is applied to the groups defined by the GROUP BY clause.
- The order of operations performed is FROM, WHERE, GROUP BY, HAVING and then ORDER BY.
- On the above example, the logic of the process will be:
  - All rows where mark is NULL are retrieved. (due to the WHERE clause)
  - The retrieved rows then are grouped into different unit\_code.
  - If the number of rows in a group is greater than 1, the unit\_code and the total is displayed. (due to the HAVING clause)



## What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
WHERE year = 2015  
GROUP BY unit_code  
HAVING avg(mark) > 30  
ORDER BY avg(mark) DESC;
```

```
SQL> SELECT unit_code, avg(mark), count(*)
2 FROM enrolmentA
3 WHERE year = 2015
4 GROUP BY unit_code
5 HAVING avg(mark) > 30
6 ORDER BY avg(mark) DESC;
```

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----  -
FIT2004      40          2
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

#### Q4. What is the output for:

```
SELECT unit_code, studid, avg(mark)
FROM enrolmentA
GROUP BY unit_code
HAVING avg(mark) > 55
ORDER BY unit_code, studid;
```

- A. FIT2094, 111, 50
- B. FIT2004, 111, 60
- C. FIT2004, 111, 60, 222, 333
- D. FIT2004, 111, 100
- E. Will print three rows
- F. Error

```
SQL> SELECT unit_code, studid, avg(mark)
2 FROM enrolmentA
3 GROUP BY unit_code
4 HAVING avg(mark) > 55
5 ORDER BY unit_code, studid;
```

Error starting at line : 1 in command -  
 SELECT unit\_code, studid, avg(mark)  
 FROM enrolmentA  
 GROUP BY unit\_code  
 HAVING avg(mark) > 55  
 ORDER BY unit\_code, studid  
 Error at Command Line : 1 Column : 19  
 Error report -  
 SQL Error: ORA-00979: **not a GROUP BY expression**  
 00979. 00000 - "not a GROUP BY expression"  
 \*Cause:  
 \*Action:

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT stu_lname, stu_fname, avg(mark)
FROM enrolment e JOIN student s
      ON s.stu_nbr = e.stu_nbr
GROUP BY s.stu_nbr;
```

The above SQL generates error message

```
SQL Error: ORA-00979: not a GROUP BY expression
00979. 00000 - "not a GROUP BY expression"
```

### **Why and how to fix this?**

- Why? Because the grouping is based on the stu\_nbr, whereas the display is based on stu\_lname and stu\_fname. The two groups may not have the same members.
- How to fix this?
  - Include the stu\_lname,stu\_fname as part of the GROUP BY condition.
- Attributes that are used in the SELECT, HAVING and ORDER BY must be included in the GROUP BY clause.

## Subqueries

- Query within a query.

"Find all students whose mark is higher than the average mark of all enrolled students"

```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg (mark)  
               FROM enrolment );
```

# Types of Subqueries

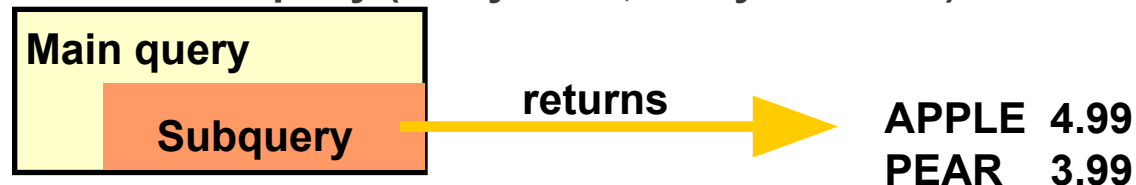
## Single-value



## Multiple-row subquery (a list of values – many rows, one column)



## Multiple-column subquery (many rows, many columns)



## Q5. What will be returned by the *inner query*?

```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg(mark)  
              FROM enrolment  
              GROUP BY unit_code);
```

- A. A value (a single column, single row).
- B. A list of values.
- C. Multiple columns, multiple rows.
- D. None of the above.



```
SQL> SELECT *  
      2 FROM enrolment  
      3 WHERE mark > (SELECT avg(mark)  
      4                FROM enrolment  
      5                GROUP BY unit_code);
```

Error starting at line : 1 in command -

SELECT \*

FROM enrolment

WHERE mark > (SELECT avg(mark)

FROM enrolment

GROUP BY unit\_code)

Error report -

**ORA-01427: single-row subquery returns more than one  
row**

## Q6. What will be returned by the *inner query*?

```
SELECT unit_code, stu_lname, stu_fname, mark
FROM enrolment e join student s
    on e.stu_nbr = s.stu_nbr
WHERE (unit_code, mark) IN (SELECT unit_code, max(mark)
    FROM enrolment
    GROUP BY unit_code);
```

- A. A value (a single column, single row).
- B. A list of values.
- C. Multiple columns, multiple rows.
- D. None of the above.

# Comparison Operators for Subquery

- Operator for single value comparison.  
=, <, >
- Operator for multiple rows or a list comparison.
  - equality
    - IN
  - inequality
    - ALL, ANY combined with <, >

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

**Q7. Which row(s) in ENROL2 table will be retrieved by the following SQL statement?**

```
SELECT * FROM enrol2
WHERE mark IN (SELECT max(mark)
               FROM enrol2
               GROUP BY unit_code);
```

- A. 1, 2, 7
- B. 7
- C. 2, 3, 7

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

```
SQL> SELECT * FROM enrol2
      2 WHERE mark IN (SELECT max(mark)
      3                     FROM enrol2
      4                     GROUP BY unit_code)
      5 ORDER BY stu_nbr, unit_code, enrol_year;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1001	2012	1	78	D
11111111	FIT1002	2013	1	80	HD
11111113	FIT1004	2013	1	89	HD

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

**Q8. Which row/s in ENROL2 will be retrieved by the following SQL statement?**

```
SELECT * FROM enrol2
WHERE mark > ANY (SELECT avg(mark)
                    FROM enrol2
                    GROUP BY unit_code);
```

- A. 1, 2, 3, 6, 7
- B. 2, 3, 7
- C. 3, 7
- D. No rows will be returned

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

```
SQL> SELECT * FROM enrol2
      2 WHERE mark > ANY (SELECT avg(mark)
      3                      FROM enrol2
      4                      GROUP BY unit_code)
      5 ORDER BY stu_nbr, unit_code, enrol_year, enrol_semester;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1001	2012	1	78	D
11111111	FIT1002	2013	1	80	HD
11111111	FIT1004	2013	1	85	HD
11111113	FIT1001	2012	2	65	C
11111113	FIT1004	2013	1	89	HD

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

**Q9. Which row/s in ENROL2 will be retrieved by the following SQL statement?**

```
SELECT * FROM enrol2
WHERE mark > ALL (SELECT avg(mark)
                    FROM enrol2
                    GROUP BY unit_code);
```

- A. 1, 2, 3, 6, 7
- B. 2, 3, 7
- C. 3, 7
- D. No rows will be returned



	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

```
SQL> SELECT * FROM enrol2
      2 WHERE mark > ALL (SELECT avg(mark)
      3                      FROM enrol2
      4                      GROUP BY unit_code)
      5 ORDER BY stu_nbr, unit_code, enrol_year, enrol_semester;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1004	2013	1	85	HD
11111113	FIT1004	2013	1	89	HD

**Q10. Find all students whose mark in any enrolled unit is lower than Wendy Wheat's lowest mark for all units she is enrolled in. What would be a possible inner query statement for the above query (assume Wendy Wheat's name is unique)?**

- A. `SELECT min(mark)  
FROM enrol2  
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`
- B. `SELECT min(mark)  
FROM enrol2 e JOIN student s on e.studid = s.studid  
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`
- C. `SELECT min(mark) FROM enrol2;`
- D. `SELECT mark  
FROM enrol2 e JOIN student s on e.studid = s.studid  
WHERE stu_lname='Wheat' AND stu_fname='Wendy';`

# Summary

- Aggregate Functions
  - count, min, max, avg, sum
- GROUP BY and HAVING clauses.
- Subquery
  - Inner vs outer query
  - comparison operators (IN, ANY, ALL)

# PART 2

## PL/SQL - Triggers (FIT3171)

# Oracle Triggers

- A trigger is PL/SQL code associated with a table, which performs an action when a row in a table is inserted, updated, or deleted.
- Triggers are used to implement some types of data integrity constraints that cannot be enforced at the DBMS design and implementation levels
- A trigger is a stored procedure/code block associated with a table
- Triggers specify a condition and an action to be taken whenever that condition occurs
- The DBMS automatically executes the trigger when the condition is met ("fires")
- A Trigger can be ENABLE'd or DISABLE'd via the ALTER command
  - ALTER TRIGGER *trigger\_name* ENABLE;

# Oracle Triggers - general form

CREATE [OR REPLACE] TRIGGER <trigger\_name>

{BEFORE | AFTER | INSTEAD OF }

{UPDATE | INSERT | DELETE}

[OF <attribute\_name>] ON <table\_name>

[FOR EACH ROW]

[WHEN]

DECLARE

BEGIN

.... *trigger body goes here* .....

END;

# Triggering Statement

BEFORE|AFTER INSERT|UPDATE [of colname]|DELETE ON Table

- The triggering statement specifies:
  - the type of SQL statement that fires the trigger body.
  - the possible options include DELETE, INSERT, and UPDATE. One, two, or all three of these options can be included in the triggering statement specification.
  - the table associated with the trigger.
- Column List for UPDATE
  - if a triggering statement specifies UPDATE, *an optional list of columns can be included in the triggering statement.*
  - if you include a column list, the trigger is fired on an UPDATE statement only when one of the specified columns is updated.
  - if you omit a column list, the trigger is fired when any column of the associated table is updated

# Trigger Body

**BEGIN**

.....

**END;**

- is a PL/SQL block that can include SQL and PL/SQL statements. These statements are executed if the triggering statement is issued and the trigger restriction (if included) evaluates to TRUE.
- Within a trigger body of a row trigger, the PL/SQL code and SQL statements have access to the **old** and **new** column values of the current row affected by the triggering statement.
- Two correlation names exist for every column of the table being modified: **one for the old column value** and **one for the new column value**.



# Correlation Names

- Oracle uses two correlation names in conjunction with every column value of the current row being affected by the triggering statement. These are denoted by:  
    OLD.ColumnName & NEW.ColumnName
  - For DELETE, only OLD.ColumnName is meaningful
  - For INSERT, only NEW.ColumnName is meaningful
  - For UPDATE, both are meaningful
- A colon must precede the OLD and NEW qualifiers when they are used in a trigger's body, but a colon is not allowed when using the qualifiers in the WHEN clause.
- Old and new values are available in both BEFORE and AFTER **row triggers**.

# FOR EACH ROW Option

- The FOR EACH ROW option determines whether the trigger is a row trigger or a statement trigger. If you specify FOR EACH ROW, the trigger fires once for each row of the table that is affected by the triggering statement. The absence of the FOR EACH ROW option means that the trigger fires only once for each applicable statement, but not separately for each row affected by the statement.

```
CREATE OR REPLACE TRIGGER display_salary_increase
AFTER UPDATE OF empmsal ON employee
FOR EACH ROW
WHEN (new.empmsal > 1000)
BEGIN
    DBMS_OUTPUT.PUT_LINE ('Employee: ' || :new.empno || ' Old salary: ' ||
        :old.empmsal || ' New salary: ' || :new.empmsal);
END;
```

# Statement Level Trigger

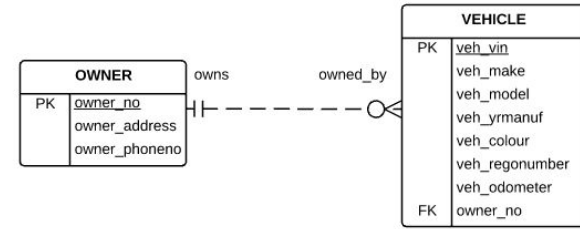
- Executed once for the whole table but will have to check all rows in the table.
- In many cases, it will be inefficient.
- **No access to the correlation values :new and :old.**

# Oracle Data FK Integrity

- Oracle offers the options:
  - UPDATE
    - no action (the default - not specified)
  - DELETE
    - no action (the default - not specified)
    - cascade
    - set null
- Subtle difference between "no action" and "restrict"
  - RESTRICT - will not allow action if child records exist, checks first
  - NO ACTION - allows action and any associated triggers, *then* checks integrity
- Databases implementations vary, for example:
  - Oracle no RESTRICT
  - IBM DB2, SQLite implement both as above



# Common use of triggers



- In the model above OWNER is the PARENT (PK end) and VEHICLE is the CHILD (FK end)
- What should the database do to maintain integrity if the user:
  - attempts to UPDATE the owner\_no of the owner (parent)
  - attempts to DELETE an owner who still has vehicles in the vehicle table
- Oracle, by default, takes the safe approach
  - UPDATE NO ACTION (no update of PK permitted if child records)
  - DELETE NO ACTION (no delete permitted if child records)
  - what if you as the developer want UPDATE CASCADE?

# Oracle Triggers

```
CREATE OR REPLACE TRIGGER Owner_Upd_Cas
BEFORE UPDATE OF owner_no ON owner
FOR EACH ROW
BEGIN
```

```
    UPDATE vehicle
    SET      owner_no = :new.owner_no
    WHERE   owner_no = :old.owner_no;
    DBMS_OUTPUT.PUT_LINE ('Corresponding owner number in the VEHICLE
table has also been updated');
END;
/
```

Implement UPDATE CASCADE rule  
OWNER 1 ---- has --- M VEHICLE  
:new.owner\_no – value of owner\_no after update  
:old.owner\_no – value of owner\_no before update

- SQL Window: To CREATE triggers, include the RUN command (/) after the last line of the file

# Common use of triggers - data integrity

- A trigger can be used to enforce user-defined integrity by triggering on a preset condition, carrying out some kind of test and then if the test fails, the trigger can raise an error (and stop the action) via a call to **raise\_application\_error**

The syntax for this call is:

```
raise_application_error(-20000, 'Error message to display');
```

the -20000 is the error number which is reported to the user, the error message is the error message the user will see. The error number can be any number less than or equal to -20000.

# Common use of triggers - data integrity - example

For example: a trigger which will ensure any unit added (ie. inserted) to the UNIT table has a unit code which starts with 'FIT'. Test your trigger and ensure it works correctly and shows your error message.

```
CREATE OR REPLACE TRIGGER check_unit_code BEFORE
    INSERT ON unit
    FOR EACH ROW
BEGIN
    IF :new.unit_code NOT LIKE 'FIT%' THEN
        raise_application_error(-20000, 'Unit code must begin with FIT');
    END IF;
END;
/
-- Test Harness
-- display before value
select * from unit;

insert into unit values ('ABC0001','Test Insert',6);

-- display after value
select * from unit;
-- closes transaction
rollback;
```



# Mutating Table

- A table that is currently being modified through an INSERT, DELETE or UPDATE statement SHOULD NOT be **read from** or **written to** because it is in a **transition state** between two stable states (before and after) where data integrity can be guaranteed.
  - Such a table is called **mutating table**.

```
CREATE OR REPLACE TRIGGER Owner_Upd_Cas BEFORE
UPDATE OF owner_no ON owner
FOR EACH ROW

DECLARE
    owner_count NUMBER;

BEGIN

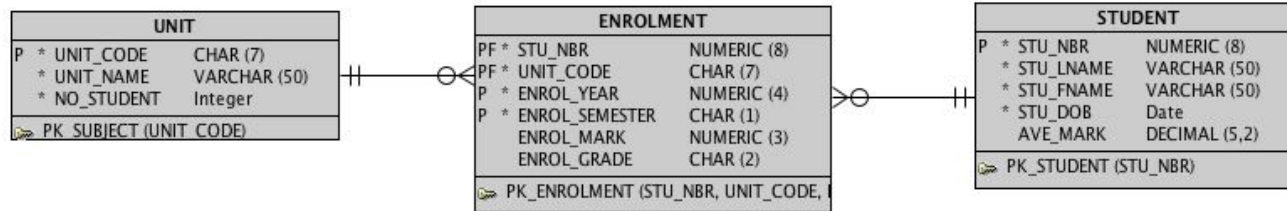
    SELECT COUNT(*) INTO owner_count
    FROM owner
    WHERE owner_no = :old.owner_no;

    IF owner_count = 1 THEN
        UPDATE vehicle
        SET owner_no = :NEW.owner_no
        WHERE owner_no = :OLD.owner_no;
        DBMS_OUTPUT.PUT_LINE ('Corresponding owner number in the VEHICLE table '
        || 'has also been updated');
    END IF;

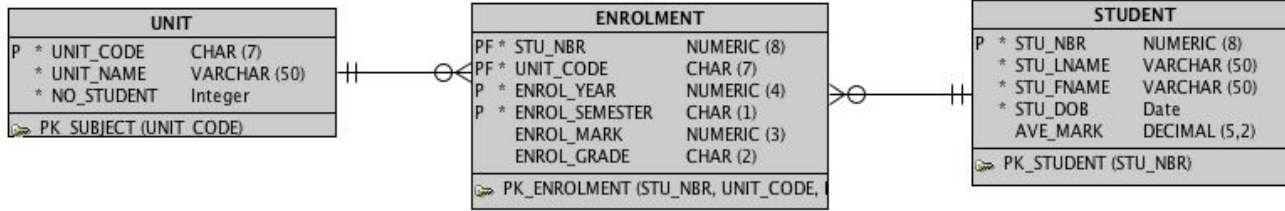
END;
```

```
update owner set owner_no = 1 where owner_no = 2
Error report -
SQL Error: ORA-04091: table LSMI1.OWNER is mutating, trigger/function may not see it
ORA-06512: at "LSMI1.OWNER_UPD_CAS", line 6
ORA-04088: error during execution of trigger 'LSMI1.OWNER_UPD_CAS'
04091. 00000 - "table %.%s is mutating, trigger/function may not see it"
*Cause:      A trigger (or a user defined plsql function that is referenced in
              this statement) attempted to look at (or modify) a table that was
              in the middle of being modified by the statement which fired it.
*Action:      Rewrite the trigger (or function) so it does not read that table.
```

# Triggers Case Study



- The student enrolment database contains two derived attributes no\_student (total number of students) and ave\_mark (average mark) .
- The total number of students is updated when **an enrolment is added or deleted**.
- The average mark is updated when **an update on attribute mark is performed**.
- For audit purpose, any deletion of enrolment needs to be recorded in an audit table. The recorded information includes the username who performed the deletion, the date and time of the deletion, the student no and unit code.



**Q5. Based on the rule to maintain the integrity of the no\_student attribute in the UNIT table as well as keeping the audit record, a trigger needs to be created for \_\_\_\_\_ table. The trigger will update a value on \_\_\_\_\_ table and insert a row to \_\_\_\_\_ table.**

- A. UNIT, ENROLMENT, AUDIT
- B. ENROLMENT, UNIT, AUDIT
- C. STUDENT, ENROLMENT, AUDIT
- D. AUDIT, UNIT, ENROLMENT

# Oracle Triggers

```
CREATE OR REPLACE TRIGGER triggername
```

```
BEFORE|AFTER  INSERT|UPDATE [of colname]|DELETE  [OR  
... ]  ON  Table
```

```
FOR EACH ROW
```

```
DECLARE
```

```
    var_name  datatype [, ...]
```

```
BEGIN
```

```
    . . . . .
```

```
END ;
```

**Q6. What would be an appropriate condition for the trigger described on the previous slide?**

- A. BEFORE INSERT OR DELETE ON enrolment.
- B. AFTER INSERT OR DELETE ON enrolment.
- C. BEFORE UPDATE OF mark ON enrolment.
- D. AFTER UPDATE OF mark ON enrolment.

```
CREATE OR REPLACE TRIGGER change_enrolment
AFTER INSERT OR DELETE ON ENROLMENT
FOR EACH ROW
DECLARE
    ???????
BEGIN
    ????????
END;
```

## Q7. What would be the logic to update the no\_student attribute in the UNIT table when a new row is inserted to ENROLMENT?

- A. UPDATE unit  
SET no\_student = no\_student + 1  
WHERE unit\_code = unit code of the inserted row
- B. UPDATE unit  
SET no\_student = (SELECT count (stu\_nbr)  
FROM enrolment  
WHERE unit\_code= unit code of the inserted row)  
WHERE unit\_code = unit code of the inserted row
- C. UPDATE unit  
SET no\_student = no\_student -1  
WHERE unit\_code = unit code of the inserted row
- D. UPDATE unit



```
CREATE OR REPLACE TRIGGER change_enrolment
AFTER INSERT OR DELETE ON ENROLMENT
FOR EACH ROW
DECLARE
    ??????
BEGIN
    IF INSERTING THEN
        UPDATE unit
        SET no_student = no_student + 1
        WHERE unit_code = :new.unit_code
    ENDIF;
    ??????
END;
```

**Q8. What would be the logic for the trigger to deal with a deletion of a row in enrolment? Assume that a table `audit_trail` contains `audit_time`, `user`, `sno` and `unitcode` attributes.**

- A. `UPDATE unit`  
    `SET no_student = no_student -1`  
    `WHERE unit_code = :old.unit_code;`
- B. `INSERT INTO audit_trail VALUES`  
    `(SYSDATE, USER,`  
        `:old.stu_nbr, :old.unit_code);`
- C. `UPDATE unit`  
    `SET no_student = no_student – 1`  
    `WHERE unit_code = :new.unit_code;`
- D. a and b.
- E. b and c.

**CREATE OR REPLACE TRIGGER change\_enrolment  
AFTER INSERT OR DELETE ON ENROLMENT  
FOR EACH ROW**

**BEGIN**

**IF INSERTING THEN**

**UPDATE unit**

**SET no\_student = no\_student + 1**

**WHERE unit\_code = :new.unit\_code;**

**END IF;**

**IF DELETING THEN**

**UPDATE unit**

**SET no\_student = no\_student -1**

**WHERE unit\_code = :old.unit\_code;**

**INSERT INTO audit\_trail VALUES (SYSDATE, USER,  
:old.stu\_nbr, :old.unit\_code);**

**END IF;**

**END;**

# Test Harness

- it is not sufficient to code a trigger only, a suitable test harness must be developed at the same time and used to ensure the trigger is working correctly.

```
-- display before value
select * from unit;

-- test the trigger for insertion
insert into enrolment values (11111111,'FIT2001',2013,2,null,null);

-- display after value
select * from unit;

-- test the trigger for deletion
delete from enrolment where stu_nbr = 11111111 and unit_code = 'FIT2001'and enrol_year =
2013 and enrol_semester = 2;

-- display after value
select * from unit; select * from audit_trail;
-- closes transaction
rollback;
```

## Statement Level Trigger

```
create or replace
TRIGGER DELETE_STATEMENT
AFTER DELETE ON ENROLMENT
BEGIN
    INSERT INTO enrol_history VALUES (SYSDATE, USER, 'Deleted');
END;
```

## Row Level Trigger

```
create or replace
TRIGGER DELETE_ENROLMENT
AFTER DELETE ON ENROLMENT
FOR EACH ROW
BEGIN
    INSERT INTO audit_trail VALUES
        (SYSDATE, USER, :old.stu_nbr, :old.unit_code);
END;
```

# Oracle Triggers

- Use triggers where:
  - a specific operation is performed, to ensure related actions are also performed
  - to enforce integrity where data has been denormalised
  - to maintain an audit trail
  - global operations should be performed, regardless of who performs the operation
  - they do NOT duplicate the functionality built into the DBMS
  - their size is reasonably small (< 50 - 60 lines of code)
- Do not create triggers where:
  - they are recursive
  - they modify or retrieve information from triggering tables